

# Short NoteBook Explanation (Author : Joel Siby)

- Socials :~
- Linkedin : <https://www.linkedin.com/in/joelsiby/>
- Github : <https://github.com/joelsiby02>
- Portfolio : <https://joelsiby.vercel.app>

## 1) Core Business Problem (and why it matters)

A growing hotel chain operates multiple bars across different locations. They are currently facing two major inventory challenges:

- **Stockouts of high-demand items (fast movers)**  
When popular brands run out, the bar cannot serve guests, which directly reduces revenue and negatively impacts guest satisfaction.
- **Overstocking of slow-moving items (slow movers)**  
When low-demand items are overstocked, it increases holding costs, ties up capital, and increases the risk of waste.

### Business goal

Build a system that helps managers decide:

**How much of each brand should each bar keep, and when should they reorder?**

## 2. Assumptions Made (and Why)?

### 1) Daily aggregation of demand

The raw data is transaction-level (timestamp included). I aggregated consumption into daily demand per **Bar** × **Alcohol Type** × **Brand** because forecasting and inventory planning require consistent daily time steps.

### 2) Missing days treated as zero demand

If a bar-brand combination had no record on a day, I assumed demand was **0** for that day. This is necessary to model intermittent demand correctly and avoid biasing forecasts.

### 3) Lead time assumed as 2 days

The dataset does not provide supplier delivery lead time. I assumed a realistic **2-day lead time** to simulate ordering and replenishment delays.

### 4) Service level fixed at 95% ( $z = 1.65$ )

I assumed a **95% service level** to balance stockout risk and holding cost. This reflects a common business tradeoff for hospitality inventory.

### 5) Bottle ordering constraint (750 ml)

Since alcohol is typically replenished in bottles, I assumed orders must be placed in **750ml multiples**. This makes the simulation more realistic and explains why inventory sometimes exceeds the Par level.

### 6) Stable operating conditions

The model assumes normal bar operations and does not explicitly account for rare events like closures, strikes, or one-time large events. These would require external signals such as hotel occupancy or event calendars.

### 3) Model Used (and why)

#### Forecasting goal

Forecast daily item-level demand (**Consumed (ml)**) for each **Bar × Alcohol Type × Brand** time series.

#### Models tested (baseline forecasting)

I tested two simple and explainable baseline forecasting approaches:

##### 1) MA7 (Rolling 7-day Moving Average)

Forecast = average demand of the last 7 days

##### 2) Seasonal Naive (7-day lag)

Forecast = demand from the same day last week

#### Why MA7 was chosen

Using a time-based split and MAE evaluation across all 96 bar-brand series:

- **MA7 performed better in 69 series (~72%)**
- **Seasonal naive performed better in 27 series (~28%)**

This indicates demand is noisy and intermittent, so MA7 is a stronger default baseline.

#### Why not more complex models?

I did not use ARIMA/Prophet/XGBoost/Deep Learning because:

- the dataset is relatively small (~6.5k rows)
- demand is intermittent with many zero-demand days
- the goal is a practical, explainable system rather than heavy model complexity

A simple baseline model is more suitable for this real-world simulation.

## 4) System Performance (and what I would improve)

### Forecasting performance

To validate forecasting quality, I compared MA7 vs Seasonal Naive using a time-based split and **MAE**.

Across 96 time series (6 bars × 16 brands):

- **MA7 performed better in 69 series (~72%)**
- Seasonal Naive performed better in 27 series (~28%)

This suggests demand is more noisy/intermittent than strongly weekly-seasonal.

---

### Inventory + simulation performance

Forecasting was converted into inventory decisions using:

- **ROP (Reorder Point)** = when to reorder
- **Par level** = target inventory after replenishment
- Lead time = 2 days
- Service level = 95%
- Orders rounded to 750ml bottles

A simulation was run to mimic real operations:

- inventory decreases daily due to demand
- orders trigger when inventory falls below ROP
- stock arrives after lead time
- stockout days and order counts are tracked

Example outcome (Brown's Bar – Yellow Tail):

- **Orders placed:** 39
  - **Stockout days:** 5
  - Overstock peaks occurred mainly due to bottle-size rounding.
-

## **Improvements (next steps)**

If more business data/time were available, I would like to improve:

**1) Use different service levels by item type**

Higher service level for fast movers, lower for slow movers.

**2) Include external demand drivers**

Weekday, holidays, hotel occupancy, events, seasonality.

**3) Better ordering optimization**

Reduce overstock caused by bottle rounding by adding caps or smarter order sizing.

**4) Use real lead times**

Learn lead time per supplier/location instead of assuming a constant value

## 5) How this solution would work in a real hotel

This system can be used as a daily decision-support tool for bar managers.

### Step-by-step workflow

- 1) **Daily update** Each bar records daily consumption and inventory balances (opening, purchases, closing).
- 2) **Demand forecasting** The system forecasts next-day demand for each **Bar × Brand** using the MA7 method.
- 3) **Inventory recommendation** For each bar-brand, the system calculates:
  - **ROP (Reorder Point)** → when to reorder
  - **Par level** → how much inventory to maintain after replenishment
- 4) **Order suggestion** If current inventory falls below ROP, the system recommends an order quantity. Orders are rounded to 750ml bottles to reflect real purchasing constraints.
- 5) **Operational simulation / monitoring** The system tracks:
  - stockout days (unmet demand)
  - order frequency
  - overstock risk (inventory above par)

### What managers get

For each bar location, managers receive:

- forecasted demand per brand
- recommended ROP and Par levels
- suggested reorder quantities
- simple KPIs (stockout risk, overstock risk)

This directly helps reduce both stockouts and overstock across all bar locations.

## 6. What would break at scale? What would I track in production?

### What could break at scale

If this system is deployed across many hotels, bars, and items, the main challenges would be:

- **Too many item-location combinations** More bars and brands means many more time series to forecast and simulate.
  - **Changing demand patterns** Demand can shift due to seasonality, new menus, promotions, or guest behavior.
  - **Unstable lead times** Supplier delays or partial deliveries can break the lead-time assumption.
  - **External events** Holidays, events or occupancy spikes can cause demand spikes that MA7 cannot predict without extra signals.
  - **Data quality issues** Missing inventory logs, incorrect consumption entries, or delayed updates will reduce reliability.
- 

### What I would track in production

To ensure the system is working correctly, I would monitor:

- **Stockout rate** (% of days with stockout)
- **Total unmet demand (ml)**
- **Average inventory level** (proxy for holding cost / overstock)
- **Inventory turnover**
- **Forecast error (MAE / bias)**
- **Order frequency and order quantity**
- **Supplier fill rate and actual lead time**